

Polynomial regression

The example belows uses a temperature-energy dataset in order to illustrate how to perform a non linear regression.

Workflow:

1. Preparation

- Extract the dataset from the `_power_demand_vstemperature.csv`
- Explore the dataset and check for missing values
- Plot the distribution
- Divide the dataset into train and test
- Create an evaluation function

2. First experiment

- Create a linear model
- Train the model on `X_train` and `y_train`
- Evaluate the model on `X_test` and `y_test`
- Visualize the prediction of the model

3. Second experiment

- Create a polynomial regression model with degree 2
- Train the model on `X_train` and `y_train`
- Evaluate the model on `X_test` and `y_test`
- Visualize the prediction of the model

4. Third experiment

- repeat the steps done in the second experiment but with degree 3

5. Third experiment

- repeat the steps done in the second experiment but with degree 4

6. Compare the evaluation of each model

```
In [ ]: # Code source: Filippo Orazi  
        # License: BSD 3 clause
```

```
import matplotlib.pyplot as plt
import numpy as np
import scipy.stats
from sklearn import datasets, linear_model
from sklearn.metrics import mean_squared_error, r2_score
import pandas as pd
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression

random_state = 42 # this will be used to guarantee the repeatability of the experiment
```

Dataset preparation

Load the dataset from a .csv file

This cell allows full compatibility between execution in Google Colab and in local

```
In [ ]: try:
        import google.colab.files
        IN_COLAB = True
    except:
        IN_COLAB = False
    # from google.colab import files
    if IN_COLAB:
        uploaded = files.upload()
```

The file must be available in the same directory, or uploaded in the Colab environment in the execution of the previous cell

Set the date column as index

```
In [ ]:
```

Out[]: **demand temp**

date		
2015-01-01	1.736065	1.7
2015-01-02	1.831672	2.2
2015-01-03	1.714934	14.4
2015-01-04	1.628577	15.6
2015-01-05	2.045394	0.0

Explore the dataset and check for missing values

In []:

Out[]:

	demand	temp
count	1096.000000	1096.000000
mean	1.831796	16.927737
std	0.329434	10.791581
min	1.316033	-15.000000
25%	1.581654	8.900000
50%	1.731479	18.900000
75%	2.024869	26.100000
max	2.804025	33.900000

In []:

The dataframe has 0 invalid rows

Create X and y


In []:

```
X has shape(1096, 1)
```

```
Y has shape(1096,)
```

Plot the distribution

```
In [ ]:
```



Divide the dataset in train and test splits

```
In [ ]:
```

Training set and test set have 767 and 329 elements respectively

Create an evaluation function to compute, print and return the metrics: rmse r2 f-statistic and p-value

```
In [ ]:
```

```
# Computation of F-statistic and p-value for the regression  
# http://facweb.cs.depaul.edu/sjost/csc423/documents/f-test-reg.htm
```

First experiment

Create a linear model

Train the model on `X_train` and `y_train`


Evaluate the model on `X_test` and `y_test`

In []:

```
Mean squared error: 0.10016
r2 score:          0.1803
f-statistic:       53.273
p-value:           2.2197e-12
```

Visualize the prediction of the model

In []:



Second experiment - Polynomial regression

We can clearly see that the linear regression model cannot really approximate the data distribution.

We can now try with a non linear regression model:

1. Use the sklearn function `PolynomialFeature` to create a new array of features. Set `degree=2` and `_includebias=False`

2. Train a Linear regression model with the new features
3. Evaluate the model
4. Visualize the predicted values of the model

Polynomial degree = 2


In []:

Out[]: `LinearRegression()`

In []:

```
Mean squared error:    0.033456
r2 score:              0.72619
f-statistic:           384.89
p-value:               1.1102e-16
```

In []:



Third experiment

Polynomial degree = 3

In []:


Polynomial degree = 3

Out[]: LinearRegression()

In []:

Mean squared error: 0.021749
r2 score: 0.822
f-statistic: 502.32
p-value: 1.1102e-16

In []:



Fourth experiment

Polynomial degree = 4

In []:

Polynomial degree = 4

Out[]: LinearRegression()


In []:

```

Mean squared error: 0.021334
r2 score:          0.8254
f-statistic:       390.05
p-value:           1.1102e-16

```

In []:



Compare the performance of the four models

In []:

Out[]:

	linear	polynomial d = 2	polynomial d = 3	polynomial d = 4
rmse	1.001591e-01	3.345625e-02	2.174942e-02	2.133387e-02
r2	-5.366169e+00	5.756325e-01	7.843318e-01	7.923317e-01
f-statistic	5.327309e+01	3.848865e+02	5.023183e+02	3.900454e+02
p-value	2.219669e-12	1.110223e-16	1.110223e-16	1.110223e-16

In []: